

Handling Missing and Inappropriate Data in a Dataset Aim: Demonstrate an experiment to handle missing data and inappropriate data in a Data set using Python Pandas Library for Data Preprocessing. Dataset Given: Hotel.csv CustomerID Age\_Group Rating(1-5) Hotel FoodPreference Bill NoOfPax EstimatedSalary Age\_Group 1 20-25 4 Ibis veg 1300 2 40000 20-25 2 30-35 5 LemonTree Non-Veg 2000 3 59000 30-35 3 25-30 6 RedFox Veg 1322 2 30000 25-30 4 20-25 -1 LemonTree Veg 1234 2 120000 20-25 5 35+ 3 Ibis Vegetarian 989 2 45000 35+ 6 35+ 3 Ibys Non-Veg 1909 2 122220 35+ 7 35+ 4 RedFox Vegetarian 1000 -1 21122 35+ 8 20-25 7 LemonTree Veg 2999 -10 345673 20-25 9 25-30 2 Ibis Non-Veg 3456 3 -99999 25-30 9 25-30 2 Ibys Non-Veg 3456 3 -99999 25-30 10 30-35 5 RedFox non-Veg - 6755 4 87777 30-35 About Dataset: No.of Columns =9 (called as series CustomerID, Age\_Group, Rating(1-5),Hotel, FoodPreference, Bill, NoOfPax, EstimatedSalary) CutomerID: Numerical Continuous data Age: Categorical Data Rating (1-5): Numerical Discrete Data Hotel: Categorical Data Food: Categorical Data Bill: Numerical Continuous data NoOfPax: Numerical Discrete EstimatedSalary: Numerical Continuous data

```
In [2]: import pandas as pd
import numpy as np

data = {
    'CustomerID': [1,2,3,4,5,6,7,8,9,9,10],
    'Age_Group': ['20-25','30-35','25-30','20-25','35+','35+','35+','20-25','25-30','25-30','30-35'],
    'Rating(1-5)': [4,5,6,-1,3,3,4,7,2,2,5],
    'Hotel': ['Ibis','LemonTree','RedFox','LemonTree','Ibis','Ibys','RedFox','LemonTree','Ibis','Ibis','RedFox'],
    'FoodPreference': ['veg','Non-Veg','Veg','Veg','Vegetarian','Non-Veg','Vegetarian','Veg','Non-Veg','Non-Veg','non-Veg'],
    'Bill': [1300,2000,1322,1234,989,1909,1000,2999,3456,3456,'-'],
    'NoOfPax': [2,3,2,2,2,2,-1,-10,3,3,4],
    'EstimatedSalary': [40000,59000,30000,120000,45000,122220,21122,345673,-99999,-99999,87777]
}

df = pd.DataFrame(data)
df['Bill'] = df['Bill'].astype(str).replace('-', np.nan)
df['Bill'] = df['Bill'].astype(float)
df['Bill'] = df['Bill'].fillna(df['Bill'].mean())
df.loc[(df['Rating(1-5)'] > 5) | (df['Rating(1-5)'] < 1), 'Rating(1-5)'] = df['Rating(1-5)'].median()
df.loc[df['NoOfPax'] <= 0, 'NoOfPax'] = df['NoOfPax'].median()
df.loc[df['EstimatedSalary'] <= 0, 'EstimatedSalary'] = df['EstimatedSalary'].median()
df['Hotel'] = df['Hotel'].replace({'Ibys': 'Ibis'})
df['FoodPreference'] = df['FoodPreference'].replace({'Vegetarian': 'Veg', 'non-Veg': 'Non-Veg', 'veg': 'Veg'})

df.drop_duplicates(inplace=True)
```

```
print(df)
print(df.info())
```

```
CustomerID  Age_Group  Rating(1-5)  Hotel  FoodPreference  Bill  \
0           1      20-25          4    Ibis        Veg  1300.0
1           2      30-35          5  LemonTree  Non-Veg  2000.0
2           3      25-30          4   RedFox        Veg  1322.0
3           4      20-25          4  LemonTree        Veg  1234.0
4           5       35+          3    Ibis        Veg   989.0
5           6       35+          3    Ibis  Non-Veg  1909.0
6           7       35+          4   RedFox        Veg  1000.0
7           8      20-25          4  LemonTree        Veg  2999.0
8           9      25-30          2    Ibis  Non-Veg  3456.0
10          10      30-35          5   RedFox  Non-Veg  1966.5
```

```
NoOfPax  EstimatedSalary
0         2            40000
1         3            59000
2         2            30000
3         2           120000
4         2            45000
5         2          122220
6         2            21122
7         2           345673
8         3            45000
10        4            87777
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 10 entries, 0 to 10
```

```
Data columns (total 8 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	CustomerID	10 non-null	int64
1	Age_Group	10 non-null	object
2	Rating(1-5)	10 non-null	int64
3	Hotel	10 non-null	object
4	FoodPreference	10 non-null	object
5	Bill	10 non-null	float64
6	NoOfPax	10 non-null	int64
7	EstimatedSalary	10 non-null	int64

```
dtypes: float64(1), int64(4), object(3)
```

```
memory usage: 720.0+ bytes
```

```
None
```

```
In [ ]: Experiment to understand the data preprocessing in Data science Description:  
Understand the importance of Data preprocessing in data science
```

```
In [3]: import pandas as pd  
from sklearn.preprocessing import MinMaxScaler, StandardScaler  
data = {  
    'CustomerID': [1, 2, 3, 4, 5],  
    'Age': [22, 25, 47, 52, 46],  
    'EstimatedSalary': [20000, 35000, 80000, 110000, 150000]  
  
}  
df = pd.DataFrame(data)  
print("Original Data:\n", df)  
scaler_minmax = MinMaxScaler()  
df_minmax = pd.DataFrame(scaler_minmax.fit_transform(df[['Age',  
    'EstimatedSalary']])),  
columns=['Age', 'EstimatedSalary'])  
print("\nAfter Min-Max Scaling:\n", df_minmax)  
scaler_std = StandardScaler()  
df_std = pd.DataFrame(scaler_std.fit_transform(df[['Age',  
    'EstimatedSalary']])),  
columns=['Age', 'EstimatedSalary'])  
print("\nAfter Standardization:\n", df_std)
```

Original Data:

	CustomerID	Age	EstimatedSalary
0	1	22	20000
1	2	25	35000
2	3	47	80000
3	4	52	110000
4	5	46	150000

After Min-Max Scaling:

	Age	EstimatedSalary
0	0.000000	0.000000
1	0.100000	0.115385
2	0.833333	0.461538
3	1.000000	0.692308
4	0.800000	1.000000

After Standardization:

	Age	EstimatedSalary
0	-1.325688	-1.234537
1	-1.083184	-0.920671
2	0.695178	0.020924
3	1.099351	0.648655
4	0.614343	1.485629

In [ ]: Experiment to understand EDA-Quantitative **and** Qualitative analysis. Description:  
Understand the importance of EDA-Quantitative **and** Qualitative analysis.

```
In [6]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Create sample dataset dictionary
data = {
    'CustomerID': [1,2,3,4,5,6,7,8,9,10],
    'Age': [22,25,47,52,46,30,28,35,40,50],
    'Gender': ['Male','Female','Female','Male','Male','Female','Female','Male','Male','Female'],
    'EstimatedSalary': [20000,35000,80000,110000,150000,60000,75000,90000,120000,50000],
    'Purchased': ['No','Yes','Yes','Yes','No','No','Yes','No','Yes','No']
}
```

```
# Convert dictionary to pandas DataFrame
df = pd.DataFrame(data)

# Display summary statistics for numerical columns
print("\nSummary Statistics:\n", df.describe())

# Display correlation matrix for numerical features
print("\nCorrelation Matrix:\n", df.corr(numeric_only=True))

# Show the distribution counts for Gender column
print("\nGender Distribution:\n", df['Gender'].value_counts())

# Show the distribution counts for Purchased column
print("\nPurchased Distribution:\n", df['Purchased'].value_counts())

# Create a figure with 4 subplots for visualization
plt.figure(figsize=(14,10))

# Plot histogram with KDE for Age distribution
plt.subplot(2,2,1)
sns.histplot(df['Age'], bins=6, kde=True, color='skyblue')
plt.title('Age Distribution')

# Plot boxplot for Estimated Salary distribution
plt.subplot(2,2,2)
sns.boxplot(x=df['EstimatedSalary'], color='lightgreen')
plt.title('Estimated Salary Boxplot')

# Plot countplot for Gender without palette to avoid warnings
plt.subplot(2,2,3)
sns.countplot(x='Gender', data=df)
plt.title('Gender Count')

# Plot countplot for Purchased without palette to avoid warnings
plt.subplot(2,2,4)
sns.countplot(x='Purchased', data=df)
plt.title('Purchase Decision Count')

# Adjust subplot spacing
plt.tight_layout()
plt.show()
```

```

# Plot correlation heatmap of numerical features
plt.figure(figsize=(6,4))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='YlGnBu')
plt.title('Correlation Heatmap')
plt.show()

# Plot pairplot showing pairwise relationships colored by 'Purchased' status
sns.pairplot(df, hue='Purchased', diag_kind='kde', palette='husl')
plt.suptitle('Pairplot - Age vs Salary vs Purchase', y=1.02)
plt.show()

```

Summary Statistics:

	CustomerID	Age	EstimatedSalary
count	10.00000	10.000000	10.000000
mean	5.50000	37.500000	79000.000000
std	3.02765	10.977249	40055.517029
min	1.00000	22.000000	20000.000000
25%	3.25000	28.500000	52500.000000
50%	5.50000	37.500000	77500.000000
75%	7.75000	46.750000	105000.000000
max	10.00000	52.000000	150000.000000

Correlation Matrix:


	CustomerID	Age	EstimatedSalary
CustomerID	1.00000	0.349361	0.329831
Age	0.349361	1.000000	0.611529
EstimatedSalary	0.329831	0.611529	1.000000


Gender Distribution:

	Gender
Male	5
Female	5
Name: count, dtype: int64	


Purchase Distribution:


	Purchased
No	5
Yes	5
Name: count, dtype: int64	





Pairplot - Age vs Salary vs Purchase





In [ ]: Experiment to understand Linear Regression **for** a given data set. Description:  
Understand the Linear regression **for** the dataset given.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
data = {
    'Experience': [1,2,3,4,5,6,7,8,9,10],
    'Salary': [25000,28000,35000,40000,45000,52000,60000,67000,75000,82000]
}
df = pd.DataFrame(data)
print("Original Data:\n", df)
X = df[['Experience']]
y = df['Salary']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("\nPredicted vs Actual:")
result = pd.DataFrame({
    'Experience': X_test.values.flatten(),
    'Actual Salary': y_test,
    'Predicted Salary': y_pred
})
print(result)

print("\nR2 Score:", r2_score(y_test, y_pred))
```

```
plt.scatter(X, y, color='blue')
plt.plot(X, model.predict(X), color='red')
plt.title('Linear Regression - Experience vs Salary')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.show()
```

Original Data:

	Experience	Salary
0	1	25000
1	2	28000
2	3	35000
3	4	40000
4	5	45000
5	6	52000
6	7	60000
7	8	67000
8	9	75000
9	10	82000

Predicted vs Actual:

	Experience	Actual Salary	Predicted Salary
2	3	35000	34653.620352
8	9	75000	73107.632094

R2 Score: 0.9953737060596429



```
In [ ]: Experiment to understand KNN algorithm for a given data set Description:  
Understand the KNN algorithm for the dataset given.
```

```
In [10]: import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import confusion_matrix, accuracy_score  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
data = {
```

```
'Age': [22, 25, 47, 52, 46, 30, 28, 35, 40, 50],  
'Salary': [20000, 35000, 80000, 110000, 150000, 60000, 75000, 90000, 120000, 50000],  
'Purchased': [0, 0, 1, 1, 1, 0, 1, 0, 1, 0]  
}  
  
df = pd.DataFrame(data)  
print("Dataset:\n", df)  
  
X = df[['Age', 'Salary']]  
y = df['Purchased']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)  
  
scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)  
X_test = scaler.transform(X_test)  
  
knn = KNeighborsClassifier(n_neighbors=3)  
knn.fit(X_train, y_train)  
  
y_pred = knn.predict(X_test)  
  
print("\nActual vs Predicted:\n")  
result = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})  
print(result)  
  
acc = accuracy_score(y_test, y_pred)  
print("\nAccuracy:", acc)  
  
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')  
plt.title('Confusion Matrix for KNN')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.show()
```


Dataset:

	Age	Salary	Purchased
0	22	20000	0
1	25	35000	0
2	47	80000	1
3	52	110000	1
4	46	150000	1
5	30	60000	0
6	28	75000	1
7	35	90000	0
8	40	120000	1
9	50	50000	0

Actual vs Predicted:

	Actual	Predicted
2	1	0
8	1	1
4	1	1

Accuracy: 0.6666666666666666



```
In [ ]: Experiment to understand Logistic Regression for a given data set. Description:  
Understand the Logistic Regression algorithm for the dataset given.
```

```
In [11]: import pandas as pd  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import confusion_matrix, accuracy_score  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
data = {
```

```
'Age': [22,25,47,52,46,30,28,35,40,50],  
'Salary': [20000,35000,80000,110000,150000,60000,75000,90000,120000,50000],  
'Purchased': [0,0,1,1,1,0,1,0,1,0]  
}  
  
df = pd.DataFrame(data)  
print("Dataset:\n", df)  
  
X = df[['Age', 'Salary']]  
y = df['Purchased']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)  
  
sc = StandardScaler()  
X_train = sc.fit_transform(X_train)  
X_test = sc.transform(X_test)  
  
model = LogisticRegression()  
model.fit(X_train, y_train)  
  
y_pred = model.predict(X_test)  
  
print("\nActual vs Predicted:")  
print(pd.DataFrame({'Actual': y_test, 'Predicted': y_pred}))  
  
acc = accuracy_score(y_test, y_pred)  
print("\nAccuracy:", acc)  
  
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, cmap='Greens', fmt='d')  
plt.title('Confusion Matrix - Logistic Regression')  
plt.xlabel('Predicted')  
plt.ylabel('Actual')  
plt.show()
```


Dataset:

	Age	Salary	Purchased
0	22	20000	0
1	25	35000	0
2	47	80000	1
3	52	110000	1
4	46	150000	1
5	30	60000	0
6	28	75000	1
7	35	90000	0
8	40	120000	1
9	50	50000	0

Actual vs Predicted:

	Actual	Predicted
2	1	0
8	1	1
4	1	1

Accuracy: 0.6666666666666666



```
In [ ]: No:20 Experiment to understand K-means clustering algorithm for a given data set.  
Description: Understand the K-means clustering algorithm for the dataset given.
```

```
In [14]: import os  
os.environ["OMP_NUM_THREADS"] = "1"  
  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
from sklearn.cluster import KMeans  
  
data = {
```

```

        'Age': [22, 25, 47, 52, 46, 30, 28, 35, 40, 50, 48, 33, 26, 45, 55],
        'Salary': [20000, 35000, 80000, 110000, 150000, 60000, 75000, 90000, 120000,
                    50000, 85000, 70000, 40000, 95000, 130000]
    }

df = pd.DataFrame(data)
print("Dataset:\n", df)

sns.scatterplot(x='Age', y='Salary', data=df, s=80, color='blue')
plt.title("Age vs Salary Distribution")
plt.show()

wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=0)
    kmeans.fit(df)
    wcss.append(kmeans.inertia_)

plt.plot(range(1, 11), wcss, marker='o', color='purple')
plt.title("Elbow Method")
plt.xlabel("Number of Clusters (K)")
plt.ylabel("WCSS")
plt.show()

kmeans = KMeans(n_clusters=3, init='k-means++', random_state=0)
df['Cluster'] = kmeans.fit_predict(df)
print("\nClustered Data:\n", df)

plt.figure(figsize=(8,6))
sns.scatterplot(x='Age', y='Salary', hue='Cluster', data=df, palette='Set2', s=100)
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],
            s=200, c='black', marker='X', label='Centroids')
plt.title("K-Means Clustering")
plt.xlabel("Age")
plt.ylabel("Salary")
plt.legend()
plt.show()

sns.pairplot(df, hue='Cluster', palette='husl', diag_kind='kde')

```


```
plt.suptitle("Pairplot of Clusters", y=1.02)
plt.show()
```

Dataset:

	Age	Salary
0	22	20000
1	25	35000
2	47	80000
3	52	110000
4	46	150000
5	30	60000
6	28	75000
7	35	90000
8	40	120000
9	50	50000
10	48	85000
11	33	70000
12	26	40000
13	45	95000
14	55	130000



```
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(  
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.
```





```
C:\Users\gurup\anaconda3\Lib\site-packages\sklearn\cluster\_kmeans.py:1419: UserWarning: KMeans is known to have a memory leak  
on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OM  
P_NUM_THREADS=1.  
    warnings.warn(
```

Clustered Data:

	Age	Salary	Cluster
0	22	20000	1
1	25	35000	1
2	47	80000	2
3	52	110000	0
4	46	150000	0
5	30	60000	1
6	28	75000	2
7	35	90000	2
8	40	120000	0
9	50	50000	1
10	48	85000	2
11	33	70000	2
12	26	40000	1
13	45	95000	2
14	55	130000	0

### K-Means Clustering





In [ ]: