

1) Write a Python program to collect, load, and perform initial exploration of the Diabetes dataset using Pandas. Display the first few records of the dataset and summarize its structure and contents.

Exp No: 5 - Data Collection and Initial Exploration

Objective: To collect, load, and perform initial exploration of the diabetes dataset


```
In [6]: # Exp No: 5 - Data Collection and Initial Exploration
# Objective: To collect, Load, and perform initial exploration of the diabetes

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Step 1: Create the diabetes.csv file manually
data = {
    "Pregnancies": [6,1,8,1,0,5,3,10,2,8],
    "Glucose": [148,85,183,89,137,116,78,115,197,125],
    "BloodPressure": [72,66,64,66,40,74,50,0,70,96],
    "SkinThickness": [35,29,0,23,35,0,32,0,45,0],
    "Insulin": [0,0,0,94,168,0,88,0,543,0],
    "BMI": [33.6,26.6,23.3,28.1,43.1,25.6,31.0,35.3,30.5,0.0],
    "DiabetesPedigreeFunction": [0.627,0.351,0.672,0.167,2.288,0.201,0.248,0.13],
    "Age": [50,31,32,21,33,30,26,29,53,54],
    "Outcome": [1,0,1,0,1,0,1,0,1,1]
}

df = pd.DataFrame(data)
df.to_csv("diabetes.csv", index=False)
print("diabetes.csv file created successfully!")

# Step 2: Load the dataset from the created CSV file
diabetes_df = pd.read_csv("diabetes.csv")

# Step 3: Display the first few rows
print("\nFirst 5 rows of the dataset:")
print(diabetes_df.head())

# Step 4: Display basic information about the dataset
print("\nDataset Information:")
print(diabetes_df.info())

# Step 5: Display statistical summary
print("\nStatistical Summary:")
print(diabetes_df.describe())

# Step 6: Check for missing values
print("\nMissing Values in Each Column:")
print(diabetes_df.isnull().sum())

# Step 7: Visualize the distribution of the target variable (Outcome)
plt.figure(figsize=(6,4))
sns.countplot(x='Outcome', data=diabetes_df)
plt.title('Distribution of Diabetes Outcome')
plt.xlabel('Outcome (0 = Non-Diabetic, 1 = Diabetic)')
plt.ylabel('Count')
plt.show()

# Step 8: Display correlation heatmap
plt.figure(figsize=(10,8))
sns.heatmap(diabetes_df.corr(), annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap of Features')
plt.show()
```


diabetes.csv file created successfully!

First 5 rows of the dataset:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Dataset Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10 entries, 0 to 9

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	10 non-null	int64
1	Glucose	10 non-null	int64
2	BloodPressure	10 non-null	int64
3	SkinThickness	10 non-null	int64
4	Insulin	10 non-null	int64
5	BMI	10 non-null	float64
6	DiabetesPedigreeFunction	10 non-null	float64
7	Age	10 non-null	int64
8	Outcome	10 non-null	int64

dtypes: float64(2), int64(7)

memory usage: 848.0 bytes

None

Statistical Summary:


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	10.00000	10.00000	10.00000	10.00000	10.00000	
mean	4.40000	127.30000	59.80000	19.90000	89.30000	
std	3.50238	40.075068	25.654976	17.978073	169.937276	
min	0.00000	78.00000	0.00000	0.00000	0.00000	
25%	1.25000	95.50000	53.50000	0.00000	0.00000	
50%	4.00000	120.50000	66.00000	26.00000	0.00000	
75%	7.50000	145.25000	71.50000	34.25000	92.50000	
max	10.00000	197.00000	96.00000	45.00000	543.00000	


	BMI	DiabetesPedigreeFunction	Age	Outcome
count	10.00000	10.00000	10.00000	10.00000
mean	27.71000	0.507800	35.90000	0.600000
std	11.259016	0.654114	11.873874	0.516398
min	0.00000	0.134000	21.00000	0.000000
25%	25.85000	0.175500	29.25000	0.000000
50%	29.30000	0.240000	31.50000	1.000000
75%	32.95000	0.558000	45.75000	1.000000
max	43.10000	2.288000	54.00000	1.000000

Missing Values in Each Column:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0

dtype: int64





2. Write a Python program to perform Time Series Analysis on a website traffic dataset. The program should load and clean the dataset, decompose the time series to identify trends and seasonality, visualize key metrics using moving averages and seasonal plots, and detect anomalies using statistical methods.

Dataset:

A dataset containing daily website traffic data with the following columns:

- Date – Date of observation
 - Page_VIEWS – Total number of page views per day
 - Unique_Visitors – Number of unique visitors per day
 - Bounce_Rate – Percentage of visitors who leave after viewing only one page
- (You can create a sample CSV file like website_traffic.csv for practice.)

Methodology:

1. Load and Clean Data:

- o Import the dataset and parse the Date column as a datetime object.
- o Handle missing values by interpolation.

2. Time Series Decomposition:

- o Decompose the time series data of Page_VIEWS to identify trend, seasonality, and residuals.

3. Visualization:

- o Plot the time series data for Page_VIEWS, Unique_Visitors, and

Bounce_Rate.

o Use moving averages to smooth the time series and highlight trends.

o Create seasonal plots to visualize patterns over time.

4. Anomaly Detection:

o Identify and visualize anomalies in the time series data using statistical methods such as the Z-score.


```
In [7]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from statsmodels.tsa.seasonal import seasonal_decompose
from scipy import stats

# --- Step 0: Generate Synthetic Website Traffic Data and Save as CSV ---

# Create a date range for 90 days
dates = pd.date_range(start='2023-01-01', periods=90, freq='D')

# Generate synthetic data with float arrays to allow NaN
np.random.seed(42)
page_views = (np.random.poisson(lam=500, size=90) + np.linspace(0, 100, 90)).astype(float)
unique_visitors = (page_views * np.random.uniform(0.6, 0.9, size=90)).astype(float)
bounce_rate = np.random.normal(loc=50, scale=10, size=90)

# Introduce some missing values
page_views[5] = np.nan
unique_visitors[20] = np.nan
bounce_rate[50] = np.nan

# Create dataframe
df_sample = pd.DataFrame({
    'Date': dates,
    'Page_Visits': page_views,
    'Unique_Visitors': unique_visitors,
    'Bounce_Rate': bounce_rate
})

# Save to CSV
df_sample.to_csv('website_traffic.csv', index=False)
print("Sample 'website_traffic.csv' created successfully!")

# --- Step 1: Load and Clean Data ---

# Load dataset with Date parsing
df = pd.read_csv('website_traffic.csv', parse_dates=['Date'])

# Sort by Date
df = df.sort_values('Date')

# Set Date as index
df.set_index('Date', inplace=True)

# Handle missing values with interpolation
df = df.interpolate()

print("\nFirst few rows of the dataset:")
print(df.head())

print("\nDataset Information:")
print(df.info())

# --- Step 2: Time Series Decomposition (Page_Visits) ---
```

```
decomposition = seasonal_decompose(df['Page_VIEWS'], model='additive', period=7)

plt.figure(figsize=(12, 8))
decomposition.plot()
plt.suptitle('Time Series Decomposition of Page Views', fontsize=14)
plt.show()

# --- Step 3: Visualization ---

# (a) Plot all metrics
plt.figure(figsize=(14, 6))
plt.plot(df.index, df['Page_VIEWS'], label='Page Views')
plt.plot(df.index, df['Unique_Visitors'], label='Unique Visitors')
plt.plot(df.index, df['Bounce_Rate'], label='Bounce Rate')
plt.title('Website Traffic Metrics Over Time')
plt.xlabel('Date')
plt.ylabel('Values')
plt.legend()
plt.show()

# (b) Moving Average on Page Views
df['PageViews_MA7'] = df['Page_VIEWS'].rolling(window=7).mean()

plt.figure(figsize=(14, 5))
plt.plot(df.index, df['Page_VIEWS'], label='Original')
plt.plot(df.index, df['PageViews_MA7'], label='7-Day Moving Average', linewidth=2)
plt.title('Page Views with Moving Average')
plt.legend()
plt.show()

# (c) Seasonal plot - Weekly pattern of Page Views
df['DayOfWeek'] = df.index.day_name()

plt.figure(figsize=(10, 5))
sns.boxplot(x='DayOfWeek', y='Page_VIEWS', data=df.reset_index(),
            order=['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday'])
plt.title('Seasonal Pattern by Day of Week')
plt.xticks(rotation=45)
plt.show()

# --- Step 4: Anomaly Detection using Z-score ---

df['Z_Score'] = np.abs(stats.zscore(df['Page_VIEWS']))

# Define anomalies as points with Z-score > 3
anomalies = df[df['Z_Score'] > 3]

plt.figure(figsize=(14, 5))
plt.plot(df.index, df['Page_VIEWS'], label='Page Views')
plt.scatter(anomalies.index, anomalies['Page_VIEWS'], color='red', label='Anomalies')
plt.title('Anomaly Detection in Page Views')
plt.legend()
plt.show()

print("\nDetected Anomalies:")
print(anomalies[['Page_VIEWS', 'Z_Score']])
```

Sample 'website_traffic.csv' created successfully!


First few rows of the dataset:




	Page_VIEWS	Unique_Visitors	Bounce_Rate
Date			
2023-01-01	492.000000	295.947090	42.696334
2023-01-02	517.123596	335.221449	52.164586
2023-01-03	476.247191	364.148192	50.455718
2023-01-04	509.370787	411.351832	43.483997
2023-01-05	528.494382	420.463988	71.439441


Dataset Information:

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 90 entries, 2023-01-01 to 2023-03-31
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Page_VIEWS      90 non-null     float64
 1   Unique_Visitors 90 non-null     float64
 2   Bounce_Rate     90 non-null     float64
dtypes: float64(3)
memory usage: 2.8 KB
None
```

<Figure size 1200x800 with 0 Axes>







Detected Anomalies:

Empty DataFrame

Columns: [Page_VIEWS, Z_Score]

Index: []

3) Random Sampling and Sampling Distribution

To explore random sampling from a population and understand the concept of sampling distribution using Python in Jupyter Notebook.

Steps:

1. Generate a Population:

- o Create a population of data with a specified distribution (e.g., normal distribution).

2. Random Sampling:

- o Perform random sampling from the population to create multiple samples of different sizes.

- o Compute sample statistics (mean, standard deviation, etc.) for each sample.

3. Sampling Distribution:

- o Plot histograms or density plots of sample statistics (e.g., sample means).
- o Compare the sampling distribution of the sample statistic (mean) with the population distribution.

4. Central Limit Theorem (Optional):

- o Demonstrate the Central Limit Theorem by showing that as sample size increases, the sampling distribution of the sample mean approaches a normal distribution regardless of the population distribution.


```
In [8]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Step 1: Generate a Population
np.random.seed(42) # For reproducibility

population_size = 100000
# Population from a normal distribution with mean=50, std=15
population = np.random.normal(loc=50, scale=15, size=population_size)

plt.figure(figsize=(10, 5))
sns.histplot(population, bins=50, kde=True, color='skyblue')
plt.title("Population Distribution (Normal Distribution)")
plt.xlabel("Value")
plt.ylabel("Frequency")
plt.show()

# Step 2 & 3: Random Sampling and Sampling Distribution
sample_sizes = [5, 30, 100]
num_samples = 1000 # Number of samples to draw for each sample size

plt.figure(figsize=(15, 12))

for i, size in enumerate(sample_sizes, 1):
    sample_means = []

    for _ in range(num_samples):
        sample = np.random.choice(population, size=size, replace=False)
        sample_means.append(np.mean(sample))

    plt.subplot(3, 1, i)
    sns.histplot(sample_means, bins=30, kde=True, color='coral')
    plt.title(f"Sampling Distribution of the Sample Mean (Sample Size = {size})")
    plt.xlabel("Sample Mean")
    plt.ylabel("Frequency")
    plt.axvline(np.mean(population), color='blue', linestyle='--', label='Population Mean')
    plt.legend()

plt.tight_layout()
plt.show()

# Step 4: Central Limit Theorem Demonstration with Non-Normal Population
# Generate a population with a uniform distribution (non-normal)
uniform_population = np.random.uniform(low=0, high=100, size=population_size)

plt.figure(figsize=(10, 5))
sns.histplot(uniform_population, bins=50, kde=True, color='lightgreen')
plt.title("Population Distribution (Uniform Distribution)")
plt.xlabel("Value")
plt.ylabel("Frequency")
plt.show()

plt.figure(figsize=(15, 12))

for i, size in enumerate(sample_sizes, 1):
    sample_means = []
```

```
for _ in range(num_samples):
    sample = np.random.choice(uniform_population, size=size, replace=False)
    sample_means.append(np.mean(sample))

    plt.subplot(3, 1, i)
    sns.histplot(sample_means, bins=30, kde=True, color='purple')
    plt.title(f"Sampling Distribution of the Sample Mean (Sample Size = {size})")
    plt.xlabel("Sample Mean")
    plt.ylabel("Frequency")
    plt.axvline(np.mean(uniform_population), color='blue', linestyle='--', label='Population Mean')
    plt.legend()

plt.tight_layout()
plt.show()
```

